A Machine Learning Application that Matters

Geoff Holmes[1], Dale Fletcher[2], Peter Reutemann[3], and Martijn van Oostrum[4]

[1] University of Waikato, Hamilton, NZ, `geoff@waikato.ac.nz`
[2] University of Waikato, Hamilton, NZ, `dale@waikato.ac.nz`
[3] University of Waikato, Hamilton, NZ, `fracpete@waikato.ac.nz`
[4] BLGG AgroXpertus, Wageningen, NL,
`martijn.vanoostrum@blgg.agroxpertus.com`

Abstract. TODO

1 Introduction

- Machine learning that matters?? See Wagstaffs paper (citations)
  http://scholar.google.com/scholar?oi=bibs&hl=en&cites=2774690908240883628
- Martijn - business side

1.1 BLGG AgroXpertus

BLGG AgroXpertus, founded in 1928, is based in Wageningen, the Netherlands, and has 200 employees. The company is part of the BLGG Group, which focuses on various types of analysis. BLGG AgroXpertus focuses on agricultural analyses. The company analyses some 500,000 samples of soil, roughage, manure, compost, nutrient solutions and crops. 'Wageningen' is a worldwide name in the fields of leading agricultural knowledge and innovations.

1.2 Efficiency improvement for farmers

The mission of BLGG AgroXpertus is to improve the efficiency of individual farmers. This is possible with innovative techniques that lead to new analytic and agricultural knowledge.

In over 80 years BLGG AgroXpertus developed a wealth of knowledge and skills in professional sampling and analysis of soil, but also of roughage, compost and manure.

BLGG AgroXpertus focuses on the areas:

- Soil nutrition and crop fertilization
- Soil Health
- Plant Health
- Roughage analysis

For the markets:

- Horticultural farming
- Arable farming
- Dairy farming
- Greenhouse growing

Due to its knowledge BLGG AgroXpertus can provide farmers worldwide with additional information, which they use to improve the yield and quality of their agricultural production, and reduce costs.

1.3 last year

- forage 140000 samples with 115000 in wageningen and 25000 in samplinq
- if you would have done all of these through ref analyses at 200 euro a samples total saving will be 28 million euros a year
- soil on average 85000 samples a year also at least 200 euro when doing all analyses though wet chemistry so 17 million euros a year..
- in a derogation season soil 150000 samples $\rightarrow$ 30 million in savings a year
- but forage is known for nir predictive analyses, soil is not

2 ADAMS

ADAMS, the Advanced Data mining and Machine learning System, is a modular, scientifc workflow engine written in Java. Currently available modules include support for Weka, MOA, R, image processing (ImageJ, JAI, ImageMagick, Gnuplot), PDF management and display, spreadsheet manipulation (CSV, Gnumeric, Excel, ODF), scripting (Groovy, Jython), GIS support (OpenStreetMap), Twitter, time-series analysis, network support (SSH, SCP, SFTP/FTP, Email), XML/HTML/JSON processing facilities and webservice capabilities.

In contrast to current versions of other workflow engines, like Kepler, KNIME or RapidMiner, where the user places operators on a canvas and connects them manually, ADAMS organizes the operators (or actors) of a workflow automatically in a tree structure without explicit connections. Instead, so called control actors determine how data flows between actors. Examples of control actors are, e.g., Sequence, Branch, Tee, Trigger, IfThenElse and Switch. Apart from flow control, there are two further aspects to an actor: functional (primitive actor or manages nested actors) and procedural in terms of input/output of data (standalone, source, transformer, sink).

Using a tree layout has advantages and disadvantages. In terms of advantages, the layout is very compact, it scales to thousands of actors, avoids having to manually rearrange the workflow in order to include additional actors (disconnect/reconnect), is context aware when adding actors (data types of input and output limit what actors can be inserted) and has customizable rules for suggesting actors depending on context for common sequences of actors. Disadvantages are, the tree layout is less intuitive compared to a canvas-based approach and it only supports 1-to-n connections. The 1-to-n limitation is mitigated using callable actors (multiple actors can channel data into a single actor using its name),

containers for storing multiple outputs, variables and internal storage (re-using data in multiple locations). Variables can be either used in expressions, e.g., ones for evaluating mathematical formulas, or attached to parameters of operators. The latter allows for influencing the flow execution, e.g., for the turning on/off of sub-flows or dynamically changing the setup of a learning algorithm. The scope of variables and internal storage can be limited using the LocalScope control actor.

Some further feature highlights: Though ADAMS is a data-driven workflow, transporting the data in so called tokens, by design rather than an event-based one, i.e., actors get executed if there is data available for them to process, it is also possible to trigger sub-flows using cronjobs. This allows, for instance, for recurring clean-up operations. Interactive actors are very useful for developing workflow applications. These actors either prompt the user to enter or select a value, controlling sub-flow execution, or require the user to inspect data, e.g., visual inspection of images, influencing the data flow. By supporting scripting (Groovy/Jython), it is possible to quickly prototype new actors without the need of compiling Java code and restarting the workflow application. Once a workflow has been developed, it is not necessary to use the graphical user interface for executing it, the command-line can be used as well (e.g., for use in a headless server environment). Java-code generation from existing flows is possible as well.

TODO screenshots, simple example

3 AgroXpertus

ADAMS - meta-flows, S2000

Features of NIR-focused ADAMS

- DataFusion: NIR and results from other tests
- Preprocessing
- Evaluation
- Outlier detection
- (automatic) retraining
- Weka (modelling) → deployment in business process
- before deployment: test multiple models in parallel on real dataset
- easy to develop (and deploy) new algorithms using Weka (open-source plugin framework)
- meta-flows: you only configure processing/prediction, actual worker flows get generated on the fly

4 Discussion

- Cost-benefit analysis (Martijn)
- Number of predictions (Martijn)
- Percentages going to wet chemistry (Martijn)

– Six Impact Challenges:
  2. $100M saved through improved decision making provided by an ML system → S2000
– Lack of Follow-through
  ADAMS → easily integrate ML system into business processes (opposed to plain Weka)

4.1 NIR

Introduction The regression method of choice for most practitioners working with near infrared spectroscopy (NIR) datasets is PLS regression. This method is fast, accurate, straightforward to implement and requires the tuning of a single parameter to maximize performance. Typically it is applied to relatively small datasets, those with fewer than 50 samples. In this study we introduce some known and some novel methods for handling NIR data and compare their performance against PLS regression for a range of NIR datasets of differing size.

Materials and Methods All data in this study originates from a FOSS 5000 instrument. The raw data is down sampled to 170 values per sample and smoothed using a Savitzky-Golay filter with a window of size 15. Two validation studies are undertaken both employing 10x10 cross-validation. The estimated RMSEP, is obtained as an average over 100 runs. In each run 90% of the data is used for training a model and 10% is held-out for testing. The corrected re-sampled t-test is then used to perform pair wise comparison between methods testing for significant differences. The first study varies the number of components (from 5 to 65) for PLS regression (PLSR) in order to find the optimum performance for that method on each dataset. The second study then carries forward these results and compares them with some well-known and novel regression methods. These other methods use sensible default, rather than optimized values for their parameters. The methods compared with PLSR are locally weighted PLS regression using 20 PLS components (LWPR), a variant of locally weighted learning that uses the PLS components to find nearest neighbors and selects the original data to build a linear regression model (LWPL), a random regression forest using PLS transformed data using 20 components (RRFP), and a number of techniques using untransformed data namely, Gaussian Processes (GP), Model trees (MT), and Support Vector Machine Regression (SVMR).

Results and Discussion Table 1 shows the number of PLS components that gave the best results in terms of RMSEP. It is interesting to note that as the size of the dataset increases so too does the number of components needed to get the best results. Table 2 summarizes the findings of comparing the best PLS regression results against the other methods. As can be seen, PLS regression is only competitive on the smallest dataset. Gaussian

Processes and LWPL with default parameters are clearly superior on larger datasets. The advantages gained by Gaussian Processes and LWPL in terms of RMSEP come at a cost. Model sizes for GP are quadratic in the number of samples, if predictions as well as prediction intervals are to be computed, or linear, if prediction intervals are not required; model sizes for LWPL are always linear in the number of samples. In an operational setting where hundreds of models are needed in memory or where models need frequent training, this can be a significant barrier to their deployment.

Dataset	Size (# of samples)	Optimum # PLS components
LACTIC	255	6
STORIG	414	39
SS	895	49
OMD	1010	62
DCAD	2522	57
K	6363	63
N	7500	63

Fig. 1. Comparison of RMSEP from varying the number of PLS components

Dataset	Best method(s) no sig diff	Worse methods i.e. sig diff
LACTIC	PLSR(6), MT, LWPL, RRFP, SVMR	GP, LWPR
STORIG	GP, LWPL	PLSR(39), MT, LWPR, RRFP, SVMR
SS	GP, LWPL	PLSR(49), MT, LWPR, RRFP, SVMR
OMD	GP, LWPL, LWPR, RRFP	PLSR(62), MT, SVMR
DCAD	GP	PLSR(57), MT, LWPR, RRFP, SVMR, LWPL
K	GP, LWPL	PLSR(63), MT, LWPR, RRFP, SVMR
N	GP	PLSR(63), MT, LWPR, RRFP, SVMR, LWPL

Fig. 2. Comparison of all methods

5 Conclusion

ADAMS matters!

References

1. Kiri L. Wagstaff (2012): Machine Learning that Matters. Proceedings of the Twenty-Ninth International Conference on Machine Learning (ICML), p. 529-536, 2012.
2. Peter Reutemann and Joaquin Vanschoren (2012): Scientific Workflow Management with ADAMS. Proceedings of the Machine Learning and Knowledge Discovery in Databases (ECML-PKDD), Part II, LNCS 7524, 2012, pp 833-837, Bristol, UK, 2012.